

Mean token usage and cost per run, by model

Wave-2 benchmark · Objectives and Script levels · 30 runs per model, 330 total · means include failed runs, which still consume tokens

