


Mean Token Consumption and Generation Cost per Run, by Model

Anthropic Google OpenAI



n = 30 runs/model, 330 total (L3/L4 only; includes failed runs, which still consume tokens). Source: benchmark_runs_wave2.csv.