


Mean Token Consumption and Generation Cost per Run, by Model (L1 only)

Anthropic Google OpenAI



n = 5 runs/model, 55 total (L1 outline generation only; 0 of 55 failed). Source: benchmark_runs_wave2.csv.