

Mean token usage and cost per run, by model (Outline level only)

Wave-2 benchmark · Outline level · 5 runs per model, 55 total · means include failed runs

Costs recomputed at current published rates (see [tracking/pricing.py](#))

