



Mean Token Consumption and Generation Cost per Run, by Model

Outline prompt level — outline generation only, not a full specification

Anthropic Google OpenAI



n = 5 runs/model, 55 total; 0 of 55 failed. Source: benchmark_runs_wave2_repriced.csv. Costs recomputed at current published rates (tracking/pricing.py).