



Mean Token Consumption and Generation Cost per Run, by Model

Objectives and Script prompt levels — full specification generation

Anthropic Google OpenAI



n = 30 runs/model, 330 total; includes failed runs, which still consume tokens. Source: benchmark_runs_wave2_repriced.csv. Costs recomputed at current published rates (tracking/pricing.py).