

Generation time and output throughput per request, by model

Wave-2 benchmark, 2026-07-10 · 30 L3/L4 runs per model · successful runs only
Time is measured end to end and includes every schema-repair retry, so it is time to a valid spec,
not single-call latency · runs executed serially, so wall-clock time is free of concurrency contention

