

Generation time by prompt specificity and target specification

Wave-2 benchmark, 2026-07-10 · successful runs · boxes pool all 11 models, whiskers to $1.5 \times \text{IQR}$, outliers omitted
L1 appears only in the left panel: it exists solely under json_lab, so pooling it into the spec comparison would bias it

