

Generation time is linear in output tokens, at a rate set by the model

Wave-2 benchmark, 2026-07-10 · successful Objectives/Script runs · shared axes, so a steeper line is a slower model
The per-model slope is what licenses summarising a model's speed as one tokens-per-second number

