

Generation time is linear in output tokens, at a rate set by the model

Wave-2 benchmark, 2026-07-10 · successful L3/L4 runs · shared axes, so a steeper line is a slower model

The per-model slope is what licenses summarising a model's speed as one tokens-per-second number

