

Token usage and cost per generated lab, by model

Wave-2 benchmark, 2026-07-10 · 35 runs per model across 7 prompt-level × spec cells · means over successful runs only

