

Total token usage and cost of the wave-2 sweep, by model

Wave-2 benchmark, 2026-07-10 · 35 runs per model across 7 prompt-level × spec cells · totals include failed runs
Costs recomputed at current published rates (see [tracking/pricing.py](#))

